Resume & LinkedIn Optimizer

This repository contains all assets for the **Agentic AI Resume Optimization Lab**. Learners will use these components to complete the Optimization Report.

1. Master Keyword Bank

A categorized list of 30+ high-signal, industry-relevant terms grouped under:

- **Orchestration**
- **Evaluation**
- **Governance**

These keywords help learners rewrite their resume bullets using modern Agentic AI terminology.

2. Target Job Description

A fictional **200-word** posting for a **Senior AI Systems Architect** at a Tier-1 tech firm.

The JD emphasizes:

- Autonomous workflows
- Production-scale RAG
- Multi-agent orchestration
- Deterministic guardrails
- Cost–latency trade-offs

Learners use this JD to align their resume bullets with role expectations.

3. Cheat Sheet

A **1-page guide** containing **5 Before & After resume bullet transformations**.

This helps learners:

- Understand how to integrate high-signal keywords
- Strengthen their resume language
- Model the transformation expected in their final submission

Master Keyword Bank (CB2 AI Roles)

Master Keyword Bank (CB2 AI Roles)

Orchestration

- Multi-Agent Orchestration
- Agentic Workflow Design
- Reflection Loop
- Critic Agent
- Tool Routing
- State Management
- Deterministic Output Parsing
- Hierarchical Reasoning
- Directed Acyclic Graphs (DAGs)
- Context Persistence

Evaluation

- Latency Optimization
- Token Efficiency
- RAGAS Evaluation
- Hallucination Reduction
- Failure Mode Analysis
- Cost–Performance Trade-off
- Model Selection Strategy
- Evaluation Pipeline Automation
- Drift Monitoring
- Benchmarking Frameworks

Governance

- Deterministic Guardrails
- Safety Constraints
- Human-in-the-Loop (HITL)
- Compliance & Governance
- Access Control Policies
- Audit Logging
- Secure Tool Invocation
- Data Privacy Controls
- Risk Mitigation Strategy
- Ethical AI Standards

Target Job Description

Target Job Description (Fictional — 200 Words)

Senior AI Systems Architect — Tier‑1 Technology Firm

Our company is seeking a Senior AI Systems Architect to design and deploy next‑generation autonomous AI workflows across mission‑critical business functions. In this role, you will architect multi‑agent systems capable of planning, reasoning, and executing complex tasks with high reliability. You will lead the design of production‑scale RAG pipelines, implement deterministic guardrails, and optimize the cost–latency trade‑offs required for enterprise deployment.

You will collaborate with cross‑functional engineering teams to build agentic workflows that integrate tool routing, state management, and secure backend operations. A key responsibility is evaluating model performance using metrics such as RAGAS, latency profiling, and hallucination reduction benchmarks. You will also design governance frameworks that ensure compliance, auditability, and safe model behavior in regulated environments.

The ideal candidate has experience with multi‑agent orchestration, LLM‑driven planning loops, and production‑grade MLOps. You must be able to translate architectural decisions into clear documentation for both technical and executive audiences. This role requires strong judgment in balancing autonomy, security, and operational cost.

If you are passionate about Software 3.0, autonomous workflows, and building scalable AI systems that deliver measurable business impact, we want to hear from you.

Cheat Sheet

Cheat Sheet: Before & After Resume Transformations

1. Before:
Built a chatbot for customer support.
After:

Architected a multi-agent reflection loop that reduced hallucination rates by 85%.

2. Before:
Improved model performance using RAG.
After:

Designed a production-scale RAG pipeline with deterministic guardrails, increasing retrieval precision by 22%.

3. Before:
Optimized LLM prompts.
After:

Implemented hierarchical reasoning with a supervisor agent, improving task completion rates by 30%.

4. Before:
Worked on model evaluation.
After:

Built an automated evaluation framework using RAGAS and latency profiling, reducing failure modes by 40%.

5. Before:
Ensured system security.
After:

Developed a governance layer with deterministic guardrails and audit logging to prevent unauthorized tool calls.